

Be proactive, take responsibility, think from the user's perspective, make thoughtful decisions, and finish what you start.

Ownership: "I took responsibility," "I didn't wait to be asked," "I owned it end-to-end," "I fixed the root cause."

Dive Deep: "I traced the issue," "I reproduced the bug," "I checked logs/metrics," "I isolated the bottleneck," "I found the root cause."

Deliver Results: "I prioritized the critical path," "I removed blockers," "I delivered within the deadline," "the result was X% improvement."

Bias for Action: "I had incomplete information," "I made a reversible decision," "I validated it quickly," "I took a calculated risk."

Learn and Be Curious: "I hadn't used this before," "I learned it quickly," "I built a prototype," "I applied it directly to the project."

Customer Obsession: "user pain point," "reduced friction," "improved the experience," "worked backward from the user."

Invent and Simplify: "automated the process," "removed unnecessary steps," "simplified the workflow," "made it reusable."

Earn Trust: "I listened to their concern," "I acknowledged my mistake," "I communicated clearly," "I used data instead of opinion."

Disagree and Commit: "I respectfully disagreed," "I presented evidence," "we discussed tradeoffs," "once we decided, I fully supported it."

Highest Standards: "temporary workaround," "permanent fix," "edge cases," "prevented recurrence," "improved reliability."

Are Right, A Lot: "I validated my assumptions," "considered alternatives," "changed my mind based on evidence."

Think Big: "designed for future scale," "made it reusable," "looked beyond the immediate requirement."

Frugality: "used existing resources," "reduced infrastructure," "optimized CPU/memory," "avoided unnecessary cost."

#	LP / Category	Interview Question	Answer Topic / Story Heading
1	Ownership	Tell me about a time you took ownership of something that wasn't explicitly your responsibility.	Code Journey – Detecting Copy/Paste Behavior
2	Ownership	Tell me about a time you noticed a problem and fixed it without being asked.	FileKit – Improving Compression with LZ77 + Huffman
3	Dive Deep	Tell me about a difficult bug or production issue you investigated.	DeepFake Analyzer – Large Image Upload Bottleneck
4	Dive Deep	Tell me about a time when data/metrics did not match what you expected.	Fake News Research – Data Cleaning + CNN-BiLSTM
5	Deliver Results	Tell me about a project where you faced major obstacles but still delivered.	EdFlix – HLS Streaming to Reduce Buffering
6	Deliver Results	Tell me about a time you had a difficult deadline.	Code Journey – Ship Core Features First
7	Failure	Tell me about your biggest failure.	Fake News Research – Re-architecting Before Fixing Data
8	Failure / Earn Trust	Tell me about a mistake you made and how you handled it.	Code Journey – False Positive Detection Design
9	Failure	Tell me about a time you made the wrong technical decision.	Image Analyzer – Monolith → Microservices
10	Learn and Be Curious	Tell me about a time you had to learn a new technology quickly.	FileKit – Learning Electron.js, Android dev & React native
11	Bias for Action	Tell me about a time you had incomplete information but had to make a decision.	SAM2-UNet – Converting 3D MRI to 2D Slices

#	LP / Category	Interview Question	Answer Topic / Story Heading
12	Customer Obsession	Tell me about a time you improved something for a customer/user.	JNTUH GPA Calculator – Simplifying Grade Calculation
13	Invent and Simplify	Tell me about something you automated or simplified.	Attendance Management – Automated Alerts
14	Earn Trust	Tell me about a conflict with a teammate.	EdFlix – Direct Streaming vs HLS
15	Have Backbone; Disagree and Commit	Tell me about a time you disagreed with your manager/team.	Image Analyzer – Monolith vs Microservices
16	Highest Standards	Tell me about a defect/problem you fixed permanently rather than applying a temporary solution.	ITTAcademy – Async Background Processing
17	Are Right, A Lot	Tell me about a difficult technical decision you made.	DeepFake Analyzer – Lossless Compression + Chunking
18	Feedback	Tell me about critical feedback you received and what you changed.	Fake News Research – Model-First → Data-First Approach
19	Ownership	Tell me about a time you had to deliver something with very little guidance.	OS Wizard – Designing Simulations from Vague Requirements
20	Dive Deep	Tell me about a problem where the obvious solution turned out to be wrong.	Image Analyzer – Optimize Payload Instead of Adding Servers
21	Deliver Results	Tell me about a time you failed to meet a goal. What happened?	GARUDA – Regex Parser vs Proper Compiler Parser
22	Learn and Be Curious	Tell me about something you taught yourself that helped your project.	FileKit – Electron.js for Offline Desktop App

#	LP / Category	Interview Question	Answer Topic / Story Heading
23	Learn and Be Curious	Tell me about a time you received feedback and changed your approach.	Research – Architecture vs Data Cleansing
24	Bias for Action	Tell me about a calculated risk you took.	SAM2-UNet – Risk of Data Loss in 3D → 2D Conversion
25	Bias for Action	Tell me about a time you had to act quickly to solve an urgent problem.	JNTUH GPA App – Floating-Point Calculation Bug
26	Customer Obsession	Tell me about a difficult customer requirement and how you handled it.	MRIT App – Offline Caching for Poor Connectivity
27	Customer Obsession	Tell me about a time you went beyond the original requirement because it improved the user experience.	MRIT App – Push Notifications + Offline Caching
28	Invent and Simplify	Tell me about a time you came up with a new solution to an existing problem.	EdFlix – Browser Device Fingerprinting
29	Earn Trust	Tell me about a time you had to convince your team to accept your idea.	Image Analyzer – Convincing Team to Use Microservices
30	Disagree and Commit	Tell me about a decision you disagreed with but supported after the decision was made.	SAM2-UNet – Supporting the 2D Model Decision
31	Frugality	Tell me about a time you achieved something despite limited resources.	ITTAcademy – Building Learning Ecosystem with Shared Infrastructure
32	Failure	Tell me about a mistake that affected other people.	Story not yet written in your file
33	Failure	Tell me about a project that didn't go as planned.	Story not yet written in your file

#	LP / Category	Interview Question	Answer Topic / Story Heading
34	Reflection	Tell me about something you would do differently if you could redo a project.	GARUDA – Design Grammar + Parser from the Beginning

Story group	Question numbers	LPs / themes
1. Code Journey / Student Code Analyzer	1, 6, 8	Ownership, Deliver Results, Failure / Earn Trust
2. FileKit	2, 10, 22	Ownership, Learn & Be Curious
3. Deepfake Image Analysis – Production Bottleneck	3, 17, 20	Dive Deep, Are Right A Lot, technical decision
4. Deepfake Image Analysis – Monolith → Microservices	9, 15, 29	Failure, Have Backbone, Earn Trust
5. Fake-News Research / Data Cleaning vs Architecture	4, 7, 18, 23	Dive Deep, Failure, Feedback, Learn & Be Curious
6. EdFlix / HLS Streaming	5, 14	Deliver Results, Earn Trust / conflict
7. SAM2-UNet / 3D → 2D MRI	11, 24, 30	Bias for Action, calculated risk, Disagree & Commit
8. JNTUH GPA Calculator	12, 25	Customer Obsession, Bias for Action
9. Attendance Automation	13	Invent & Simplify
10. OS Wizard	19	Ownership / ambiguity
11. GARUDA Compiler	21, 34	Failure, reflection / what you'd redo
12. College App / Offline Caching	26, 27	Customer Obsession
13. EdFlix Device Fingerprinting	28	Invent & Simplify
14. ITTacademy Ecosystem	31	Frugality

1. Ownership - Tell me about a time you took ownership of something that wasn't explicitly your responsibility. (Code Journey Development)

Situation: As a TA during my master's, I noticed it was hard to tell whether students were writing code themselves or pasting large sections.

Task: It wasn't part of my assigned responsibilities, but I felt it affected fair evaluation, so I decided to address it.

Action: I built a platform that tracked coding behavior like keystrokes, edits, deletions, cursor movement, and paste events. Instead of making an automatic judgment, it flagged unusual sessions for professor review.

Result: It gave the professor an additional signal to evaluate submissions fairly and taught me to take ownership of problems even when they fall outside my formal responsibilities.

2. Ownership - Tell me about a time you noticed a problem and fixed it without being asked. (FileKit – Squash compression)

Situation: I built an open-source file management system, and during testing I noticed my Huffman-based compression was reducing file sizes by only about 30%.

Task: No one asked me to improve it, but I felt the compression performance wasn't good enough, so I took ownership of optimizing it.

Action: I analyzed the bottleneck and realized Huffman handled symbol-level redundancy but not repeated byte sequences. I researched alternatives and added **LZ77 before Huffman encoding** to eliminate repeated sequences first.

Result: The average compression improved from roughly **30% to over 50%**, with even better results on some files. It reinforced my habit of proactively improving weaknesses instead of stopping once something simply works.

3. Dive Deep - Tell me about a difficult bug or production issue you investigated.

1. (DeepFake Analysis Image Upload Bottleneck)

Situation: I built an image-analysis model using DlibCNN, VGG16, and Transformers. It worked well locally, but larger images caused bottlenecks in production because we initially had only one processing instance.

Task: I needed to make large-image processing reliable without changing the ML model.

Action: I traced the issue to payload size and transfer overhead. I added lossless compression, split large images into smaller chunks, sent them independently, and reconstructed them on the server using the request identifier before inference.

Result: This removed the large-image bottleneck and made production processing more reliable. It also taught me to debug the full system, not just the model, because production issues can come from networking and infrastructure as much as the algorithm itself.

4. Dive Deep - Tell me about a time when the data/metrics did not match what you expected.

Situation: In my fake-news detection research, traditional ML models were consistently underperforming.

Task: I needed to determine whether the issue was the model or the quality of the text data.

Action: I analyzed the dataset, removed noisy and irrelevant text, normalized the inputs, and retrained after each change. Once preprocessing improved, I found the traditional models still struggled with context, so I moved to a **CNN-BiLSTM** architecture.

Result: The new approach produced much stronger evaluation metrics and became the final research model. It taught me to diagnose poor performance systematically **data first, then preprocessing, then model architecture**.

5. Deliver Results - Tell me about a project where you faced major obstacles but still delivered. (EdFlix video buffering, -> HLS protocol)

Situation:

During my master's, while working as TA, I was teaching through an online classroom platform, but as usage grew, relying only on live sessions became difficult to scale.

Task:

I decided to build **EdFlix**, an on-demand learning platform that could reliably deliver course content to users with different network conditions.

Action:

I built it using React and Spring Boot. The biggest obstacle was buffering for users on slower networks. I traced the problem to serving large video files directly and redesigned the streaming layer using **HLS, M3U8 playlists, and smaller video segments** so clients could stream content progressively instead of downloading large files.

Result:

I successfully launched the platform, it grew to **500+ daily users**, and buffering dropped by about **40%**. It taught me to identify the real system bottleneck and redesign the architecture instead of applying small fixes around the symptom.

6. Deliver Results - Tell me about a time you had a difficult deadline. (Code journey story, shipped core features first then later released UI/UX updates)

Situation: During my master's, I was balancing coursework, TA duties, grading, and two research papers while also needing to deliver a student code-analysis tool.

Task: I had to build it within a limited timeline without letting my other responsibilities slip.

Action: I broke the work into milestones, prioritized the core functionality first, keystrokes, edits, cursor movement, and paste patterns and added reporting afterward. I also blocked dedicated time for grading, coursework, and research.

Result: I delivered the tool on time while meeting my other commitments. It taught me to prioritize the highest-impact work and execute effectively under competing deadlines.

7. Failure - Tell me about your biggest failure. (Re-architecting instead of Re-cleaning data)

Situation: In my fake-news detection research, the model metrics were lower than expected, so I initially spent too much time experimenting with different architectures.

Task: I needed to improve performance enough to produce reliable research results.

Action: After several iterations, I realized the bigger issue was the dataset. I cleaned noisy samples, improved preprocessing, and then finalized the CNN-BiLSTM architecture.

Result: Performance improved, but the bigger lesson was to validate **data, assumptions, and bottlenecks before redesigning the architecture.**

8. Failure / Earn Trust - Tell me about a mistake you made and how you handled it. (Code journey initial version)

Situation: While building the student code analyzer, I realized my first approach could incorrectly flag legitimate coding behavior.

Task: That was a serious issue because a false positive could unfairly penalize a student.

Action: I redesigned the system to use multiple signals typing patterns, cursor movement, edits, timing, and paste events and changed it to only flag sessions for professor review instead of making an automatic judgment.

Result: This reduced the risk of unfair conclusions and taught me that when software can affect people, **false positives and edge cases must be designed for from the start.**

9. Failure - Tell me about a time you made the wrong technical decision.

1. (Deployment of image analyzer on monolithic arch, instead of micro-services based)

Situation:

As a Research Assistant, I deployed an image-analysis model using a monolithic architecture. It worked initially, but as concurrent requests increased, heavy inference requests started slowing down the entire application.

Task:

I realized my original architecture was limiting scalability and needed to be corrected.

Action:

I redesigned the system into microservices, separating preprocessing, inference, and supporting services so each component could scale independently.

Result:

The new design handled higher load more reliably and made scaling much more efficient. My key learning was to choose architecture based on expected workload and scaling needs, not just initial simplicity.

10. Learn and Be Curious - Tell me about a time you had to learn a new technology quickly. (Android development & React native dev for JNTUH GPA calculator / Learning electron JS, to build users trust)

Situation: I initially built FileKit as a web-based file-management application.

Task: I wanted it to work fully offline, feel native, and have better access to local system resources.

Action: I quickly learned **Electron.js**, integrated it with my Java and React stack, implemented IPC communication, and handled platform-specific behavior across Windows, Linux, and macOS.

Result: I converted FileKit into a cross-platform desktop application while preserving the core functionality. It taught me to learn new technologies based on the product need, not just stay within the stack I already know.

11. Bias for Action - Tell me about a time you had incomplete information but had to make a decision. (3D to 2D dataset, for SAM2-Unet)

Situation: In a research project, I had a **3D BraTS MRI dataset**, but the assigned SAM2-UNet model expected **2D inputs**.

Task: I needed to adapt the data without losing important tumor information.

Action: I built a preprocessing pipeline that converted the MRI volumes into 2D slices and prioritized slices with larger tumor regions for training and evaluation.

Result: This let me continue the research without changing the model or dataset. It taught me to make a sound technical decision with limited guidance, validate it quickly, and iterate based on results.

12. Customer Obsession - Tell me about a time you improved something for a customer/user.

(Made marks calculation easy for students)

Situation: During my undergraduate studies at JNTUH, I noticed students struggled with manually calculating GPA because of credits and grading rules.

Task: I wanted to make the process simple and quick.

Action: I built a **JNTUH GPA calculator** that automatically applied the university grading rules based on entered grades and credits, with a simple interface.

Result: It reduced manual effort and made GPA calculation easier for students. It taught me to identify user friction and solve it with the **simplest effective solution**.

13. Invent and Simplify - Tell me about something you automated or simplified. (Attendance management story)

Situation:

While working with students, attendance tracking and manually notifying absent students was repetitive and time-consuming.

Task:

I wanted to automate the process so instructors didn't have to continuously calculate absences and send individual reminders.

Action:

I built an attendance management system that automatically tracked each student's absences, sent an email whenever a student was absent, and monitored their overall attendance. If a student's attendance dropped below the required threshold, the system automatically notified the professor.

Result:

This reduced manual tracking and follow-up work and gave both students and professors timely visibility into attendance issues.

14. Earn Trust - Tell me about a conflict with a teammate. (Normal streaming VS HLS)

Situation: While building EdFlix, a teammate wanted to keep direct video streaming because it was simpler, but I was concerned about buffering on slower networks.

Task: We needed to decide whether the simpler approach was good enough for production.

Action: I tested both approaches under different network conditions and showed that direct streaming degraded significantly. I proposed HLS with M3U8 segments, and after discussing the tradeoffs, we agreed to switch.

Result: Buffering dropped by about 40%, and I learned to resolve technical disagreements with data rather than opinion.

15. Have Backbone; Disagree and Commit -Tell me about a time you disagreed with your manager/team. (Deployment on monolithic arch vs micro-services based)

Situation: While working as a Research Assistant, our image-analysis application was deployed as a monolith. The team preferred keeping it that way because it was simpler to maintain.

Task: As concurrent inference requests increased, I felt scaling the entire application for one compute-heavy component was inefficient.

Action: I raised the concern with the team and backed it with load-test results. I proposed separating preprocessing and model inference into independent services so we could scale only the bottleneck. After discussing the added complexity and tradeoffs, the team agreed to the change.

Result: We moved to a microservice-based architecture, which improved scalability and resource utilization. I learned that when I disagree, I should bring data and alternatives not just an opinion.

16. Highest Standards - Tell me about a defect/problem you fixed permanently rather than applying a temporary solution. (Sync -> Async)

Situation: At ITTacademy, some APIs were slow because background work like email notifications was being executed synchronously.

Task: A temporary fix would have been increasing timeouts or adding more server resources, but that wouldn't solve the root cause.

Action: I redesigned the flow by separating non-critical work from the request path and introduced an asynchronous, thread-pooled background processing pipeline.

Result: API wait time dropped by about 70%, and the system could handle background jobs much more efficiently. The key lesson was to fix the architecture causing the problem instead of masking the symptom.

17. Are Right, A Lot - Tell me about a difficult technical decision you made. (Applying lossless compression and chunking image, for deep-fake analysis)

Situation: While deploying my image-analysis model, larger images created a bottleneck in production because the system was running on a limited number of instances.

Task: I had to decide whether to simply scale the infrastructure or redesign how image data was transferred and processed.

Action: I chose to first optimize the pipeline. I applied lossless compression, split large images into smaller chunks, sent them independently, and reconstructed them on the server using the request identifier before inference.

Result: This reduced the transfer bottleneck and allowed the existing infrastructure to handle larger images more reliably. It taught me to optimize the system bottleneck before solving every performance problem by adding more infrastructure.

18. Feedback - Tell me about critical feedback you received and what you changed.

Situation: During my fake-news detection research, I was spending a lot of time experimenting with different model architectures because the metrics were below expectations.

Task: My professor gave me critical feedback that I was focusing too much on the model and not enough on understanding the dataset itself.

Action: I took the feedback seriously and changed my approach. I analyzed the data more deeply, removed noisy and irrelevant samples, improved preprocessing, and only then continued model experimentation with CNN-BiLSTM.

Result: The evaluation results improved, and I became much more systematic in how I approach technical problems. Now I validate the data and assumptions before making major architectural changes.

19. Ownership - Tell me about a time you had to deliver something with very little guidance. (OS Wizard story)

Situation:

I was asked to build **OS Wizard**, a web application covering 10 operating-system problems, but I was given almost no guidance on the UI, UX, or how the simulations should work.

Task:

I had to decide how to present complex OS computations in a way that students could easily follow.

Action:

I designed the interface myself and broke each problem into step-by-step visual simulations. I focused on showing how inputs changed, how each computation progressed, and why each result was produced instead of just displaying the final answer.

Result:

I delivered a usable learning tool that made OS concepts easier for students to understand. The experience taught me how to take a vague requirement, define the product behavior myself, and deliver without needing detailed instructions.

20. Dive Deep - Tell me about a problem where the obvious solution turned out to be wrong.

(Adding more servers VS pre-processing image & chunking)

Situation:

While deploying my image-analysis model, larger images caused significant latency in production.

Task:

The obvious solution was to add more server instances, but before scaling infrastructure, I wanted to confirm the actual bottleneck.

Action:

I traced the request flow and found that much of the delay came from transferring large image payloads rather than model inference itself. Instead of simply adding servers, I applied **lossless compression**, split larger images into chunks, and reconstructed them on the server before inference.

Result:

The pipeline handled larger images more reliably without immediately increasing infrastructure. I learned to validate the root cause before applying the most obvious solution.

21. Deliver Results - Tell me about a time you failed to meet a goal. What happened? (GARUDA story)

Situation: I started building **GARUDA**, my own language processor, with the goal of creating a full compiler pipeline.

Task: I wanted to build parsing, memory management, and instruction translation from Java into C.

Action: To move quickly, I initially used regex-based parsing. It worked for simple instructions, but became hard to maintain as the language grew. I had planned to replace it with a proper parser, but underestimated the scope.

Result: I delivered a working language processor, but not the full compiler I originally planned. It taught me to **scope ambitious projects into smaller milestones and validate the architecture early**.

22. Learn and Be Curious - Tell me about something you taught yourself that helped your project. (Android Dev, Cross platform dev / Learnt electron.js to give UI/UX for my offline apps)

Situation:

While building FileKit, I had already developed the core file-management functionality and a web interface, but I wanted the application to work fully offline and feel like a native desktop product.

Task:

I had never worked deeply with **Electron.js**, so I needed to learn it quickly and integrate it with my existing Java and React application.

Action:

I taught myself Electron, learned IPC communication between processes, and worked through platform-specific file-system behavior across Windows, Linux, and macOS.

Result:

I converted FileKit into a cross-platform desktop application with native-like offline access. It taught me how to independently learn a new technology and apply it directly to a real product need.

23. Learn and Be Curious - Tell me about a time you received feedback and changed your approach. (Research story, Arch vs Data cleansing)

Situation: During my fake-news detection research, I was spending a lot of time experimenting with different model architectures because the metrics were below expectations.

Task: My professor pointed out that I was focusing too much on the model and not enough on understanding the underlying data.

Action: I changed my approach and first analyzed the dataset, removed noisy and irrelevant samples, improved preprocessing, and then evaluated which architecture actually fit the data best.

Result: The model performance improved, and since then I've followed a **data-first, root-cause-first approach** before making major architectural changes.

24. Bias for Action - Tell me about a calculated risk you took. (Risk of losing data while converting from 3D to 2D)

Situation: I had to use a **2D SAM2-UNet model** with a **3D BraTS MRI dataset**.

Task: I needed to convert the data without losing important tumor information.

Action: I converted the volumes into 2D slices, prioritized slices with larger tumor regions, and validated the approach experimentally before full training.

Result: This let me use the assigned model without changing the dataset or architecture. It taught me to **take calculated risks, understand the downside, and validate quickly before scaling the decision**.

25. Bias for Action - Tell me about a time you had to act quickly to solve an urgent problem. (Incorrect floating-point operations on JNTUH app)

Situation: I built an Android GPA calculator for students, but after release I found that I had not handled floating-point division correctly, which caused inaccurate results in some cases.

Task: I needed to give students a correct solution immediately, but publishing an Android update would take time because of the app-store review process.

Action: I fixed the calculation logic right away and quickly deployed a web version so students could continue using the calculator correctly while the mobile update was under review.

Result: Students had an immediate working alternative, and I released the corrected Android version about a week later. It taught me that during an urgent issue, I should first restore a reliable user experience, then follow through with the permanent fix.

26. Customer Obsession - Tell me about a difficult customer requirement and how you handled it. (PDF/Data offline caching in MRIT app, for improper net connection)

Situation: Students wanted real-time college news and academic updates, but many also had inconsistent internet connectivity.

Task: I needed to make sure the app stayed useful even when users were offline or on a weak network.

Action: I combined Firebase Realtime Database and REST APIs for live updates, added push notifications for important announcements, and implemented offline caching so previously fetched content was still available without connectivity.

Result: The app served **1,200+ active users** and improved accessibility and engagement by about **80%**. It taught me to design around actual user constraints, not just ideal conditions.

27. Customer Obsession - Tell me about a time you went beyond the original requirement because it improved the user experience. (Added push notifications & messaging service)

Situation:

The original requirement for the college app was simply to provide students with news and academic updates.

Task:

While using the app, I realized students could still miss important updates when they had poor connectivity or weren't actively checking the application.

Action:

Beyond the original requirement, I added **push notifications and offline caching** so important announcements reached students proactively and previously loaded content remained available without internet access.

Result:

The app became much more reliable and convenient for students, eventually serving **1,200+ active users** and improving accessibility and engagement by about **80%**.

28. Invent and Simplify - Tell me about a time you came up with a new solution to an existing problem. (Digital finger print for registered devices for EdFlix)

Situation:

While building EdFlix subscriptions, I wanted to prevent one account from being freely shared across multiple devices.

Task:

In a browser-based application, I couldn't reliably access a device's physical hardware identifier, so I needed another way to recognize a registered device.

Action:

I created a lightweight device fingerprint using browser information, screen characteristics, and a generated unique token. I stored that association with the user account and validated it during login before allowing access.

Result:

This gave me a practical way to limit account sharing without requiring native hardware access or additional software installation. It taught me to work within platform constraints and design a simpler alternative when the ideal information isn't available.

29. Earn Trust - Tell me about a time you had to convince your team to accept your idea. (Deployment on Micro-services-based arch)

Situation:

As a Research Assistant, our image-analysis system was initially deployed as a monolith. The team preferred keeping it because it was simpler to maintain.

Task:

As inference traffic increased, I felt scaling the whole application for one compute-heavy component was inefficient, so I proposed separating the workload.

Action:

I ran load tests, identified inference as the main bottleneck, and showed the team how independent services would let us scale preprocessing and inference separately. I also explained the tradeoff: slightly more operational complexity in exchange for better scalability and isolation.

Result:

The team agreed to move toward a microservice-based design, which improved scalability and resource utilization. I learned that the best way to influence a technical decision is to bring data, tradeoffs, and a practical migration path.

30. Disagree and Commit - Tell me about a decision you disagreed with but supported after the decision was made. (Working with 3D based model instead of State-of-the-art SAM2-Unet)

Situation:

During my research, we were asked to use a **2D SAM2-Unet model** with the **3D BraTS MRI dataset**. I initially felt a 3D-based approach would preserve the volumetric information better.

Task:

I raised my concern, but the project direction was to continue with SAM2-Unet.

Action:

Once the decision was made, I committed to it and focused on making the approach work. I converted the 3D MRI volumes into 2D slices and prioritized slices containing larger tumor regions so we retained the most useful information for segmentation.

Result:

We were able to successfully use the assigned model with the BraTS dataset and complete the research work. I learned that after respectfully raising a concern, I need to fully support the agreed direction and make it successful.

31. Frugality - Tell me about a time you achieved something despite limited resources.

1. (ITTAcademy/ Messaging app with auto deletion option, as extra cost for DB)

Situation:

While building ITTAcademy, I had very limited budget and infrastructure, but I still wanted to provide students with a complete learning ecosystem.

Task:

I needed to support live classes, recorded content, and collaborative coding without relying on expensive third-party platforms.

Action:

I built the core products myself: an **online classroom**, a **GitHub-like code repository with web and CLI access**, and **EdFlix** for on-demand learning. I reused common backend components, open-source technologies, and shared infrastructure instead of purchasing separate services for each requirement.

Result:

I was able to build and operate a complete learning ecosystem with very little spend while still supporting real students and multiple use cases. It taught me to solve resource constraints through reuse and good architecture before adding cost.

32. Failure - Tell me about a mistake that affected other people.

33. Failure - Tell me about a project that didn't go as planned.

**34. Reflection - Tell me about something you would do differently if you could redo a project.
(GARUDA re-start story)**

Situation:

While building my GARUDA language processor, I initially used regular expressions to parse instructions because it allowed me to build the first version quickly.

Task:

As I added more instructions, the parsing logic became increasingly difficult to maintain and extend.

Action:

If I could redo the project, I would define the language grammar first and build a proper tokenizer and parser from the beginning instead of relying heavily on regex.

Result:

That would have required more upfront effort, but it would have made the language much easier to extend, debug, and maintain. The project taught me to think about long-term extensibility before choosing the fastest initial implementation.

Why Amazon?

- Amazon's **scale** is a major reason I'm interested
 - I want to build systems where **reliability, latency, and performance decisions** have real customer impact.
- I like the **ownership culture**,
 - where engineers are expected to **understand the problem and take responsibility** for what they build.
- The **customer obsession** mindset strongly matches how I approach engineering.
- What really stands out to me is Amazon's evolution from an e-commerce company into a platform that now powers other businesses through **AWS and cloud infrastructure**.
- I also researched the breadth of engineering at Amazon - from **Alexa+ and agentic AI**, to **Trainium/Graviton chips**, to autonomous systems like **Zoox**.
- My interests are in **backend, distributed, and performance-critical systems**, so Amazon gives me the kind of scale and technical depth I'm looking for.
- Overall, I want to work in an environment where I can **own impactful products and solve technically challenging problems at massive scale**.

Intro

- Hi, I'm Rohit.
- I completed my Master's in Computer Science from Texas Tech University.
- Working as a Research Assistant. AI research, particularly deepfake detection.
- I primarily work with Java, Spring Boot, Python, React, AWS, and Azure.
- I've built production-grade systems including video streaming platforms, real-time applications, asynchronous processing pipelines, and scalable APIs.
- I enjoy solving scalability and reliability problems and taking systems from design through deployment.
- I'm now looking for an SDE role where I can work on larger-scale systems and continue growing as an engineer.

Usage of LLMs as an SDE

1. What is an LLM?

An LLM is a neural network trained on large amounts of text to predict the next token. Through that training, it learns language patterns and can perform tasks like summarization, coding, question answering, and generation.

2. What is a token?

A token is a small unit of text processed by an LLM. It may be a whole word, part of a word, punctuation, or a symbol.

3. What is a context window?

The context window is the maximum amount of text or tokens the model can consider in one request, including the prompt, conversation history, and generated response.

4. What is hallucination?

Hallucination is when an LLM generates information that sounds reasonable but is factually incorrect or unsupported.

To reduce it, I can use:

RAG, trusted sources, structured prompts, validation, and human review for critical decisions.

5. What is RAG?

Retrieval-Augmented Generation retrieves relevant information from an external knowledge source and provides it to the LLM as context before generating the answer.

Example:

User question → embedding → retrieve relevant company documents → send documents + question to LLM → generate answer.

6. Why use RAG?

I would use RAG when the model needs access to **private, frequently changing, or domain-specific information**.

It also helps reduce hallucinations by grounding responses in retrieved content.

7. What are embeddings?

Embeddings are numerical vector representations of text that capture semantic meaning.

Texts with similar meanings generally have vectors that are closer together.

8. What is a vector database?

A vector database stores embeddings and supports similarity search.

In a RAG system, I can convert the user's question into an embedding and retrieve the most semantically similar documents.

9. RAG vs fine-tuning?

RAG changes what information the model receives.

Fine-tuning changes the model's behavior or learned patterns.

I would use RAG for changing/private knowledge and fine-tuning when I want consistent behavior, style, or task specialization.

10. How would you build an LLM system that answers questions from company documents?

I would:

Documents → clean/chunk → generate embeddings → store in vector DB → embed user query → retrieve relevant chunks → send chunks + query to LLM → return answer with sources.

That's a basic RAG architecture.

11. How would you reduce hallucinations?

I would:

- ground answers using RAG
- provide clear instructions
- restrict the model to trusted context
- ask it to say when information is unavailable
- validate important outputs
- include source citations where possible

12. How would you evaluate an LLM application?

I wouldn't evaluate only whether the response “sounds good.”

I would measure things like:

accuracy, relevance, groundedness, hallucination rate, latency, cost, and user satisfaction.

For RAG, I'd separately evaluate retrieval quality and generation quality.

13. What is temperature?

Temperature controls randomness in generation.

Lower temperature: more deterministic and consistent.

Higher temperature: more diverse and creative.

For something like code generation or factual Q&A, I would generally prefer a lower temperature.

14. How would you reduce LLM latency?

I could:

- use a smaller/faster model where appropriate
- reduce unnecessary prompt/context size
- retrieve fewer but more relevant documents
- cache repeated requests
- stream responses
- parallelize independent operations

15. How would you reduce LLM cost?

I would:

- use smaller models for simpler tasks
- reduce token usage
- cache common results
- summarize long conversation history
- retrieve only relevant context
- batch requests when appropriate
- use larger models only for harder tasks

16. How would you protect sensitive data when using an LLM?

I would avoid unnecessarily sending sensitive information to the model.

I would use:

access controls, data minimization, encryption, redaction, secure providers, logging controls, and clear retention policies.

17. What happens if the prompt exceeds the context window?

Some content must be removed, summarized, or retrieved dynamically.

For a large document system, I would generally use **chunking + retrieval** instead of sending every document to the LLM.

18. What is prompt engineering?

Prompt engineering means structuring instructions and context so the model produces more reliable output.

For example, instead of:

“Analyze this.”

I could say:

“Analyze this error log, identify the likely root cause, list supporting evidence, and suggest three fixes ordered by risk.”

19. LLM vs traditional ML?

Traditional ML is usually trained for a specific task, such as classification.

An LLM is a general-purpose language model capable of performing many tasks from instructions.

If I need a small, predictable classification system with low latency, traditional ML may still be the better choice.

20. What is an AI agent?

An LLM normally generates a response.

An **agent** can use the LLM to reason about a task and then interact with tools or APIs to take actions.

For example:

User: “Schedule a meeting.”

An agent might inspect calendars, find availability, and call a calendar API.

21. What if the LLM gives different answers to the same question?

LLMs are probabilistic, so some variation is expected.

If consistency matters, I can:

- lower temperature
- use stronger system instructions
- provide structured context
- constrain output format
- validate results programmatically

22. How would you make an LLM answer only using trusted information?

I would use RAG with trusted sources and instruct the model:

“Answer only using the provided context. If the answer isn't present, say you don't have enough information.”

For critical applications, I'd also validate citations or retrieved evidence.

23. When would you NOT use an LLM?

I wouldn't use an LLM when:

- deterministic logic is sufficient
- exact mathematical correctness is required
- latency/cost constraints are very strict
- a simple database lookup solves the problem
- incorrect output could cause serious harm without validation

A good engineer shouldn't use an LLM just because it's available.

24. What are the main limitations of LLMs?

Main limitations include:

hallucination, limited context, nondeterminism, outdated/missing knowledge, cost, latency, privacy concerns, and inability to guarantee correctness.

25. How do you use LLMs in your daily software-development work?

A good answer for **you**:

“I use LLMs mainly as a productivity multiplier. I use them to understand unfamiliar APIs, summarize documentation, generate initial test cases, review code for edge cases, and brainstorm debugging approaches. I also use them to improve technical documentation. But I don't treat their output as automatically correct I verify important technical decisions through documentation, tests, and my own reasoning.”